

AHCI

The specification: <https://www.intel.com/content/dam/www/public/us/en/documents/technical-specifications/serial-ata-ahci-spec-rev1-3-1.pdf>

Contents

Introduction

SATA basic

Find an AHCI controller

Determining what mode the controller is in

AHCI Registers and Memory Structures

Detect attached SATA devices

AHCI port memory space initialization

AHCI & ATAPI

Example - Read hard disk sectors

Checklist

Initialisation

Start read/write command

IRQ handler

External Links

Introduction

AHCI (Advance Host Controller Interface) is developed by Intel to facilitate handling SATA devices. The AHCI specification emphasizes that an AHCI controller (referred to as host bus adapter, or HBA) is designed to be a data movement engine between system memory and SATA devices. It encapsulates SATA devices and provides a standard PCI interface to the host. System designers can easily access SATA drives using system memory and memory mapped registers, without the need for manipulating the annoying task files as IDE do.

An AHCI controller may support up to 32 ports which can attach different SATA devices such as disk drives, port multipliers, or an enclosure management bridge. AHCI supports all native SATA features such as command queueing, hot plugging, power management, etc. To a software developer, an AHCI controller is just a PCI device with bus master capability.

AHCI is a new standard compared to IDE, which has been around for twenty years. There exists little documentation about its programming tips and tricks. Possibly the only available resource is the Intel AHCI specification (see External Links) and some open source operating systems such as Linux. This article shows the minimal steps an OS (not BIOS) should do to put AHCI controller into a workable state, how to identify drives attached, and how to read physical sectors from a SATA disk. To keep concise, many technical details and deep explanations of some data structures have been omitted.

It should be noted that IDE also supports SATA devices and there are still debates about which one, IDE or AHCI, is better. Some tests even show that a SATA disk acts better in IDE mode than AHCI mode. But the common idea is that AHCI performs better and will be the standard PC to SATA interface, though some driver software should be enhanced to fully cultivate AHCI capability.

All the diagrams in this article are copied from the Intel AHCI specification 1.3.

SATA basic

There are at least two SATA standards maintained respectively by T13 (<http://www.t13.org>) and SATA-IO (<http://www.sata-io.org>). The SATA-IO focuses on serial ATA and T13 encompasses traditional parallel ATA specifications as well.

While the hardware specifications for IDE and SATA (and even between different devices implementing them) differ greatly, the API and ABI are very similar. To a software developer, the biggest difference between SATA and parallel ATA is that SATA uses FIS (Frame Information Structure) packet to transport data between host and device. An FIS can be viewed as a data set of traditional task files, or an encapsulation of ATA commands. SATA uses the same command set as parallel ATA.

1) FIS types

Following code defines different kinds of FIS specified in Serial ATA Revision 3.0.

```
typedef enum
{
    FIS_TYPE_REG_H2D      = 0x27, // Register FIS - host to device
    FIS_TYPE_REG_D2H      = 0x34, // Register FIS - device to host
    FIS_TYPE_DMA_ACT      = 0x39, // DMA activate FIS - device to host
    FIS_TYPE_DMA_SETUP    = 0x41, // DMA setup FIS - bidirectional
    FIS_TYPE_DATA          = 0x46, // Data FIS - bidirectional
    FIS_TYPE_BIST          = 0x58, // BIST activate FIS - bidirectional
    FIS_TYPE_PIO_SETUP    = 0x5F, // PIO setup FIS - device to host
    FIS_TYPE_DEV_BITS     = 0xA1, // Set device bits FIS - device to host
} FIS_TYPE;
```

2) Register FIS – Host to Device

A host to device register FIS is used by the host to send command or control to a device. As illustrated in the following data structure, it contains the IDE registers such as command, LBA, device, feature, count and control. An ATA command is constructed in this structure and issued to the device. All reserved fields in an FIS should be cleared to zero.

```
typedef struct tagFIS_REG_H2D
{
    // DWORD 0
    uint8_t fis_type; // FIS_TYPE_REG_H2D

    uint8_t pmport:4; // Port multiplier
    uint8_t rsv0:3;   // Reserved
    uint8_t c:1;      // 1: Command, 0: Control

    uint8_t command; // Command register
    uint8_t feature; // Feature register, 7:0

    // DWORD 1
    uint8_t lba0;    // LBA low register, 7:0
    uint8_t lba1;    // LBA mid register, 15:8
    uint8_t lba2;    // LBA high register, 23:16
    uint8_t device;  // Device register

    // DWORD 2
    uint8_t lba3;    // LBA register, 31:24
    uint8_t lba4;    // LBA register, 39:32
    uint8_t lba5;    // LBA register, 47:40
    uint8_t featureh; // Feature register, 15:8

    // DWORD 3
    uint8_t countl;   // Count register, 7:0
    uint8_t counth;   // Count register, 15:8
    uint8_t icc;      // Isochronous command completion
    uint8_t control;  // Control register

    // DWORD 4
    uint8_t rsv1[4]; // Reserved
} FIS_REG_H2D;
```

3) Register FIS – Device to Host

A device to host register FIS is used by the device to notify the host that some ATA register has changed. It contains the updated task files such as status, error and other registers.

```
typedef struct tagFIS_REG_D2H
{
    // DWORD 0
    uint8_t fis_type; // FIS_TYPE_REG_D2H

    uint8_t pmport:4; // Port multiplier
    uint8_t rsv0:2;   // Reserved
    uint8_t i:1;      // Interrupt bit
    uint8_t rsv1:1;   // Reserved
```

```

    uint8_t status;          // Status register
    uint8_t error;          // Error register

    // DWORD 1
    uint8_t lba0;           // LBA low register, 7:0
    uint8_t lba1;           // LBA mid register, 15:8
    uint8_t lba2;           // LBA high register, 23:16
    uint8_t device;         // Device register

    // DWORD 2
    uint8_t lba3;           // LBA register, 31:24
    uint8_t lba4;           // LBA register, 39:32
    uint8_t lba5;           // LBA register, 47:40
    uint8_t rsv2;           // Reserved

    // DWORD 3
    uint8_t countl;          // Count register, 7:0
    uint8_t counth;          // Count register, 15:8
    uint8_t rsv3[2];         // Reserved

    // DWORD 4
    uint8_t rsv4[4];         // Reserved
} FIS_REG_D2H;

```

4) Data FIS – Bidirectional

This FIS is used by the host or device to send data payload. The data size can be varied.

```

typedef struct tagFIS_DATA
{
    // DWORD 0
    uint8_t fis_type; // FIS_TYPE_DATA

    uint8_t pmport:4; // Port multiplier
    uint8_t rsv0:4;   // Reserved

    uint8_t rsv1[2]; // Reserved

    // DWORD 1 ~ N
    uint32_t data[1]; // Payload
} FIS_DATA;

```

5) PIO Setup – Device to Host

This FIS is used by the device to tell the host that it's about to send or ready to receive a PIO data payload.

```

typedef struct tagFIS_PIO_SETUP
{
    // DWORD 0
    uint8_t fis_type; // FIS_TYPE_PIO_SETUP

    uint8_t pmport:4; // Port multiplier
    uint8_t rsv0:1;   // Reserved
    uint8_t d:1;      // Data transfer direction, 1 - device to host
    uint8_t i:1;      // Interrupt bit
    uint8_t rsv1:1;

    uint8_t status;    // Status register
    uint8_t error;    // Error register

    // DWORD 1
    uint8_t lba0;     // LBA low register, 7:0
    uint8_t lba1;     // LBA mid register, 15:8
    uint8_t lba2;     // LBA high register, 23:16
    uint8_t device;   // Device register

    // DWORD 2
    uint8_t lba3;     // LBA register, 31:24
    uint8_t lba4;     // LBA register, 39:32
    uint8_t lba5;     // LBA register, 47:40
    uint8_t rsv2;     // Reserved

    // DWORD 3
    uint8_t countl;   // Count register, 7:0
    uint8_t counth;   // Count register, 15:8
    uint8_t rsv3;     // Reserved
    uint8_t e_status; // New value of status register

    // DWORD 4
    uint16_t tc;      // Transfer count
}

```

```

    uint8_t rsv4[2]; // Reserved
} FIS PIO SETUP;

```

6) DMA Setup – Device to Host

```

typedef struct tagFIS_DMA_SETUP
{
    // DWORD 0
    uint8_t fis_type; // FIS_TYPE_DMA_SETUP

    uint8_t pmport:4; // Port multiplier
    uint8_t rsv0:1; // Reserved
    uint8_t d:1; // Data transfer direction, 1 - device to host
    uint8_t i:1; // Interrupt bit
    uint8_t a:1; // Auto-activate. Specifies if DMA Activate FIS is needed

    uint8_t rsvd[2]; // Reserved

    //DWORD 1&2

    uint64_t DMAbufferID; // DMA Buffer Identifier. Used to Identify DMA buffer in host memory.
                          // SATA Spec says host specific and not in Spec. Trying AHCI spec might work.

    //DWORD 3
    uint32_t rsvd; //More reserved

    //DWORD 4
    uint32_t DMAbufOffset; //Byte offset into buffer. First 2 bits must be 0

    //DWORD 5
    uint32_t TransferCount; //Number of bytes to transfer. Bit 0 must be 0

    //DWORD 6
    uint32_t resvd; //Reserved

} FIS_DMA_SETUP;

```

7) Example

This example illustrates the steps to read the Identify data from a device. Error detection and recovery is ignored.

To issue an ATA Identify command to the device, the FIS is constructed at follows.

```

FIS_REG_H2D fis;
memset(&fis, 0, sizeof(FIS_REG_H2D));
fis.fis_type = FIS_TYPE_REG_H2D;
fis.command = ATA_CMD_IDENTIFY; // 0xEC
fis.device = 0; // Master device
fis.c = 1; // Write command register

```

After the device receives this FIS and successfully read the 256 words data into its internal buffer, it sends a PIO Setup FIS – Device to Host to tell the host that it's ready to transfer data and the data size (FIS_PIO_SETUP.tc).

After the PIO Setup FIS – Device to Host has been sent correctly, the device sends a DATA FIS to the host which contains the received data payload (FIS_DATA.data).

This scenario is described in SATA revision 3.0 as a PIO data-in command protocol. But an AHCI controller will do the latter two steps for the host. The host software needs only setup and issue the command FIS, and tells the AHCI controller the memory address and size to store the received data. After everything is done, the AHCI controller will issue an interrupt (if enabled) to notify the host to check the data.

Find an AHCI controller

An AHCI controller can be found by enumerating the PCI bus. It has a class id 0x01 (mass storage device) and normally a subclass id 0x06 (serial ATA, but can be IDE Interface 0x01). The vendor id and device id should also be checked to ensure it's really an AHCI controller.

Determining what mode the controller is in

As you may be aware, a SATA controller can either be in IDE emulation mode or in AHCI mode. The problem that enters here is simple:

How to find what mode the controller is in. The documentation is really obscure on this. Perhaps the best way is to initialize a SATA controller as both IDE and AHCI. In this way, as long as you are careful about non-existent ports, you cannot go wrong.


One possible way of doing this is by checking the bit 31 of GHC register. It's labeled as AHCI Enable.

AHCI Registers and Memory Structures

As mentioned above, host communicates with the AHCI controller through system memory and memory mapped registers. The last PCI base address register (BAR[5], header offset 0x24) points to the AHCI base memory, it's called ABAR (AHCI Base Memory Register). All AHCI registers and memories can be located through ABAR. The other PCI base address registers act same as a traditional IDE controller. Some AHCI controller can be configured to simulate a legacy IDE one.

1) HBA memory registers

An AHCI controller can support up to 32 ports. HBA memory registers can be divided into two parts: Generic Host Control registers and Port Control registers. Generic Host Control registers controls the behavior of the whole controller, while each port owns its own set of Port Control registers. The actual ports an AHCI controller supported and implemented can be calculated from the Capacity register (HBA_MEM.cap) and the Port Implemented register (HBA_MEM.pi).



```
typedef volatile struct tagHBA_MEM
{
    // 0x00 - 0x2B, Generic Host Control
    uint32_t cap;           // 0x00, Host capability
    uint32_t ghc;           // 0x04, Global host control
    uint32_t is;             // 0x08, Interrupt status
    uint32_t pi;             // 0x0C, Port implemented
    uint32_t vs;             // 0x10, Version
}
```

```

uint32_t ccc_ctl; // 0x14, Command completion coalescing control
uint32_t ccc_pts; // 0x18, Command completion coalescing ports
uint32_t em_loc; // 0x1C, Enclosure management location
uint32_t em_ctl; // 0x20, Enclosure management control
uint32_t cap2; // 0x24, Host capabilities extended
uint32_t bohc; // 0x28, BIOS/OS handoff control and status

// 0x2C - 0x9F, Reserved
uint8_t rsv[0xA0-0x2C];

// 0xA0 - 0xFF, Vendor specific registers
uint8_t vendor[0x100-0xA0];

// 0x100 - 0x10FF, Port control registers
HBA_PORT ports[1]; // 1 ~ 32
} HBA_MEM;

typedef volatile struct tagHBA_PORT
{
    uint32_t clb; // 0x00, command list base address, 1K-byte aligned
    uint32_t clbu; // 0x04, command list base address upper 32 bits
    uint32_t fb; // 0x08, FIS base address, 256-byte aligned
    uint32_t fbu; // 0x0C, FIS base address upper 32 bits
    uint32_t is; // 0x10, interrupt status
    uint32_t ie; // 0x14, interrupt enable
    uint32_t cmd; // 0x18, command and status
    uint32_t rsv0; // 0x1C, Reserved
    uint32_t tfd; // 0x20, task file data
    uint32_t sig; // 0x24, signature
    uint32_t ssts; // 0x28, SATA status (SCR0:SStatus)
    uint32_t sctl; // 0x2C, SATA control (SCR2:SControl)
    uint32_t serr; // 0x30, SATA error (SCR1:SError)
    uint32_t sact; // 0x34, SATA active (SCR3:SActive)
    uint32_t ci; // 0x38, command issue
    uint32_t sntf; // 0x3C, SATA notification (SCR4:SNotification)
    uint32_t fbs; // 0x40, FIS-based switch control
    uint32_t rsv1[11]; // 0x44 ~ 0x6F, Reserved
    uint32_t vendor[4]; // 0x70 ~ 0x7F, vendor specific
} HBA_PORT;


```

This memory area should be configured as uncachable as they are memory mapped hardware registers, not normal prefetchable RAM. For the same reason, the data structures are declared as "volatile" to prevent the compiler from over optimizing the code.

2) Port Received FIS and Command List Memory

Each port can attach a single SATA device. Host sends commands to the device using Command List and device delivers information to the host using Received FIS structure. They are located at HBA_PORT.clb/clbu, and HBA_PORT.fb/fbu. The most important part of AHCI initialization is to set correctly these two pointers and the data structures they point to.

Port System Memory Structures



3) Received FIS

There are four kinds of FIS which may be sent to the host by the device as indicated in the following structure declaration. When an FIS has been copied into the host specified memory, an according bit will be set in the Port Interrupt Status register (HBA_PORT.is).

Data FIS – Device to Host is not copied to this structure. Data payload is sent and received through PRDT (Physical Region Descriptor Table) in Command List, as will be introduced later.

```

typedef volatile struct tagHBA_FIS
{
    // 0x00
    FIS_DMA_SETUP    dsfis;      // DMA Setup FIS
    uint8_t          pad0[4];

    // 0x20
    FIS_PIO_SETUP    psfis;      // PIO Setup FIS
    uint8_t          pad1[12];

    // 0x40
    FIS_REG_D2H rjis;        // Register - Device to Host FIS
    uint8_t          pad2[4];

    // 0x58
    FIS_DEV_BITS     sdbfis;     // Set Device Bit FIS

    // 0x60
    uint8_t          ufis[64];

    // 0xA0
    uint8_t          rsv[0x100-0xA0];
} HBA_FIS;

```


4) Command List

Host sends commands to the device through Command List. Command List consists of 1 to 32 command headers, each one is called a slot. Each command header describes an ATA or ATAPI command, including a Command FIS, an ATAPI command buffer and a bunch of Physical Region Descriptor Tables specifying the data payload address and size.

To send a command, the host constructs a command header, and set the according bit in the Port Command Issue register (HBA_PORT.ci). The AHCI controller will automatically send the command to the device and wait for response. If there are some errors, error bits in the Port Interrupt register (HBA_PORT.is) will be set and additional information can be retrieved from the Port Task File register (HBA_PORT.tfd), SStatus register (HBA_PORT.ssts) and SError register (HBA_PORT.serr). If it succeeds, the Command Issue register bit will be cleared and the received data payload, if any, will be copied from the device to the host memory by the AHCI controller.

How many slots a Command List holds can be got from the Host capability register (HBA_MEM.cap). It must be within 1 and 32. SATA supports queued commands to increase throughput. Unlike traditional parallel ATA drive; a SATA drive can process a new command when an old one is still running. With AHCI, a host can send up to 32 commands to device simultaneously.

Command List Structure



```

typedef struct tagHBA_CMD_HEADER
{
    // DW0
    uint8_t cfl:5;      // Command FIS length in DWORDS, 2 ~ 16
    uint8_t a:1;         // ATAPI
    uint8_t w:1;         // Write, 1: H2D, 0: D2H
    uint8_t p:1;         // Prefetchable

    uint8_t r:1;         // Reset
    uint8_t b:1;         // BIST
    uint8_t c:1;         // Clear busy upon R_OK
    uint8_t rsv0:1;      // Reserved
    uint8_t pmp:4;       // Port multiplier port

    uint16_t prdtl;     // Physical region descriptor table length in entries

    // DW1
    volatile
    uint32_t prdbc;     // Physical region descriptor byte count transferred

    // DW2, 3
}

```

```

    uint32_t ctba;      // Command table descriptor base address
    uint32_t ctbau;     // Command table descriptor base address upper 32 bits

    // DW4 - 7
    uint32_t rsv1[4];   // Reserved
} HBA_CMD_HEADER;


```

5) Command Table and Physical Region Descriptor Table

As described above, a command table contains an ATA command FIS, an ATAPI command buffer and a bunch of PRDT (Physical Region Descriptor Table) specifying the data payload address and size.

A command table may have 0 to 65535 PRDT entries. The actual PRDT entries count is set in the command header (HBA_CMD_HEADER.prdtl). As an example, if a host wants to read 100K bytes continuously from a disk, the first half to memory address A1, and the second half to address A2. It must set two PRDT entries, the first PRDT.DBA = A1, and the second PRDT.DBA = A2.

An AHCI controller acts as a PCI bus master to transfer data payload directly between device and system memory.



```

typedef struct tagHBA_CMD_TBL
{
    // 0x00
    uint8_t cfis[64]; // Command FIS

    // 0x40
    uint8_t acmd[16]; // ATAPI command, 12 or 16 bytes

    // 0x50
    uint8_t rsv[48]; // Reserved
}

```

```

// 0x80
HBA_PRDT_ENTRY prdt_entry[1]; // Physical region descriptor table entries, 0 ~ 65535
} HBA_CMD_TBL;

typedef struct tagHBA_PRDT_ENTRY
{
    uint32_t dba;           // Data base address
    uint32_t dbau;          // Data base address upper 32 bits
    uint32_t rsv0;          // Reserved

    // DW3
    uint32_t dbc:22;        // Byte count, 4M max
    uint32_t rsv1:9;         // Reserved
    uint32_t i:1;            // Interrupt on completion
} HBA_PRDT_ENTRY;

```

Detect attached SATA devices

1) Which port has a device attached

As specified in the AHCI specification, firmware (BIOS) should initialize the AHCI controller into a minimal workable state. OS usually needn't reinitialize it from the bottom. Much information is already there when the OS boots.

The Port Implemented register (HBA_MEM.pi) is a 32 bit value and each bit represents a port. If the bit is set in this register, then the corresponding port is present (note: this does not necessarily mean a device is attached to the port). A device is attached to a port if the value of the port's HBA_PxSSTS.det field is 0x3.

2) What kind of device is attached

There are four kinds of SATA devices, and their signatures are defined as below. The Port Signature register (HBA_PORT.sig) contains the device signature, just read this register to find which kind of device is attached at the port. Some buggy AHCI controllers may not set the Signature register correctly. The most reliable way is to judge from the Identify data read back from the device.

```

#define SATA_SIG_ATA     0x00000101 // SATA drive
#define SATA_SIG_ATAPI   0xEB140101 // SATAPI drive
#define SATA_SIG_SEMB    0xC33C0101 // Enclosure management bridge
#define SATA_SIG_PM      0x96690101 // Port multiplier

#define AHCI_DEV_NULL 0
#define AHCI_DEV_SATA 1
#define AHCI_DEV_SEMB 2
#define AHCI_DEV_PM 3
#define AHCI_DEV_SATAPI 4

#define HBA_PORT_IPM_ACTIVE 1
#define HBA_PORT_DET_PRESENT 3

void probe_port(HBA_MEM *abar)
{
    // Search disk in implemented ports
    uint32_t pi = abar->pi;
    int i = 0;
    while (i<32)
    {
        if (pi & 1)
        {
            int dt = check_type(&abar->ports[i]);
            if (dt == AHCI_DEV_SATA)
            {
                trace_ahci("SATA drive found at port %d\n", i);
            }
            else if (dt == AHCI_DEV_SATAPI)
            {
                trace_ahci("SATAPI drive found at port %d\n", i);
            }
            else if (dt == AHCI_DEV_SEMB)
            {
                trace_ahci("SEMB drive found at port %d\n", i);
            }
            else if (dt == AHCI_DEV_PM)
            {
                trace_ahci("PM drive found at port %d\n", i);
            }
            else
            {
                trace_ahci("No drive found at port %d\n", i);
            }
        }
        pi = pi >> 1;
        i++;
    }
}

```

```

        }

        pi >= 1;
        i++;
    }

// Check device type
static int check_type(HBA_PORT *port)
{
    uint32_t ssts = port->ssts;

    uint8_t ipm = (ssts >> 8) & 0x0F;
    uint8_t det = ssts & 0x0F;

    if (det != HBA_PORT_DET_PRESENT) // Check drive status
        return AHCI_DEV_NULL;
    if (ipm != HBA_PORT_IPM_ACTIVE)
        return AHCI_DEV_NULL;

    switch (port->sig)
    {
        case SATA_SIG_ATAPI:
            return AHCI_DEV_SATAPI;
        case SATA_SIG_SEMB:
            return AHCI_DEV_SEMB;
        case SATA_SIG_PM:
            return AHCI_DEV_PM;
        default:
            return AHCI_DEV_SATA;
    }
}

```

AHCI port memory space initialization

BIOS may have already configured all the necessary AHCI memory spaces. But the OS usually needs to reconfigure them to make them fit its requirements. It should be noted that Command List must be located at 1K aligned memory address and Received FIS be 256 bytes aligned.

Before rebasing Port memory space, OS must wait for current pending commands to finish and tell HBA to stop receiving FIS from the port. Otherwise an accidentally incoming FIS may be written into a partially configured memory area. This is done by checking and setting corresponding bits at the Port Command And Status register (HBA_PORT.cmd). The example subroutines stop_cmd() and start_cmd() do the job.

The following example assumes that the HBA has 32 ports implemented and each port contains 32 command slots, and will allocate 8 PRDTs for each command slot. (Note: unlike in the above struct definitions, this is using 8 instead of 1)

```

#define AHCI_BASE 0x4000000 // 4M

#define HBA_PxCMD_ST 0x0001
#define HBA_PxCMD_FRE 0x0010
#define HBA_PxCMD_FR 0x4000
#define HBA_PxCMD_CR 0x8000

void port_rebase(HBA_PORT *port, int portno)
{
    stop_cmd(port); // Stop command engine

    // Command list offset: 1K*portno
    // Command list entry size = 32
    // Command list entry maxim count = 32
    // Command list maxim size = 32*32 = 1K per port
    port->c1b = AHCI_BASE + (portno<<10);
    port->c1bu = 0;
    memset((void*)(port->c1b), 0, 1024);

    // FIS offset: 32K+256*portno
    // FIS entry size = 256 bytes per port
    port->fb = AHCI_BASE + (32<<10) + (portno<<8);
    port->fbu = 0;
    memset((void*)(port->fb), 0, 256);

    // Command table offset: 40K + 8K*portno
    // Command table size = 256*32 = 8K per port
    HBA_CMD_HEADER *cmdheader = (HBA_CMD_HEADER*)(port->c1b);
    for (int i=0; i<32; i++)
    {
        cmdheader[i].prdtl = 8; // 8 prdt entries per command table
        // 256 bytes per command table, 64+16+48+16*8
    }
}

```

```

// Command table offset: 40K + 8K*portno + cmdheader_index*256
cmdheader[i].ctba = AHCI_BASE + (40<<10) + (portno<<13) + (i<<8);
cmdheader[i].ctbau = 0;
memset((void*)cmdheader[i].ctba, 0, 256);
}

start_cmd(port); // Start command engine
}

// Start command engine
void start_cmd(HBA_PORT *port)
{
    // Wait until CR (bit15) is cleared
    while (port->cmd & HBA_PxCMD_CR)
        ;

    // Set FRE (bit4) and ST (bit0)
    port->cmd |= HBA_PxCMD_FRE;
    port->cmd |= HBA_PxCMD_ST;
}

// Stop command engine
void stop_cmd(HBA_PORT *port)
{
    // Clear ST (bit0)
    port->cmd &= ~HBA_PxCMD_ST;

    // Clear FRE (bit4)
    port->cmd &= ~HBA_PxCMD_FRE;

    // Wait until FR (bit14), CR (bit15) are cleared
    while(1)
    {
        if (port->cmd & HBA_PxCMD_FR)
            continue;
        if (port->cmd & HBA_PxCMD_CR)
            continue;
        break;
    }
}

```

AHCI & ATAPI

The documentation regarding using the AHCI interface to access an ATAPI device (most likely an optical drive) is rather poorly explained in the specification. However, once you understand that the HBA does most of the work for you it is rather simple. The AHCI/APATA method works by issuing the ATA PACKET command (0xA0) instead of the READ (READ is shown in the example below) and populating the ACMD field of the HBA_CMD_TBL with the 12/16 byte ATAPI command and setting the 'a' field to 1 in the HBA_CMD_HEADER which tells the HBA to perform the multi-step process (all done automatically) of transmitting the PACKET command, then sending the ATAPI device the ACMD.

Example - Read hard disk sectors

The code example shows how to read "count" sectors from sector offset "startl:starth" to "buf" with LBA48 mode from HBA "port". Every PRDT entry contains 8K bytes data payload at most.

```

#define ATA_DEV_BUSY 0x80
#define ATA_DEV_DRQ 0x08

bool read(HBA_PORT *port, uint32_t startl, uint32_t starth, uint32_t count, uint16_t *buf)
{
    port->is = (uint32_t) -1; // Clear pending interrupt bits
    int spin = 0; // Spin lock timeout counter
    int slot = find_cmds_slot(port);
    if (slot == -1)
        return false;

    HBA_CMD_HEADER *cmdheader = (HBA_CMD_HEADER*)port->clb;
    cmdheader += slot;
    cmdheader->cfl = sizeof(FIS_REG_H2D)/sizeof(uint32_t); // Command FIS size
    cmdheader->w = 0; // Read from device
    cmdheader->prdtl = (uint16_t)((count-1)>>4) + 1; // PRDT entries count

    HBA_CMD_TBL *cmdbtbl = (HBA_CMD_TBL*)(cmdheader->ctba);
    memset(cmdbtbl, 0, sizeof(HBA_CMD_TBL) +
           (cmdheader->prdtl-1)*sizeof(HBA_PRDT_ENTRY));

    // 8K bytes (16 sectors) per PRDT

```

```

for (int i=0; i<cmdheader->prdtl-1; i++)
{
    cmdtbl->prdt_entry[i].dba = (uint32_t) buf;
    cmdtbl->prdt_entry[i].dbc = 8*1024-1; // 8K bytes (this value should always be set to 1 less than the actual value)
    cmdtbl->prdt_entry[i].i = 1;
    buf += 4*1024; // 4K words
    count -= 16; // 16 sectors
}
// Last entry
cmdtbl->prdt_entry[i].dba = (uint32_t) buf;
cmdtbl->prdt_entry[i].dbc = (count<<9)-1; // 512 bytes per sector
cmdtbl->prdt_entry[i].i = 1;

// Setup command
FIS_REG_H2D *cmdfis = (FIS_REG_H2D*)(&cmdtbl->cfis);

cmdfis->fis_type = FIS_TYPE_REG_H2D;
cmdfis->c = 1; // Command
cmdfis->command = ATA_CMD_READ_DMA_EX;

cmdfis->lba0 = (uint8_t)startl;
cmdfis->lbal = (uint8_t)(startl>>8);
cmdfis->lba2 = (uint8_t)(startl>>16);
cmdfis->device = 1<<6; // LBA mode

cmdfis->lba3 = (uint8_t)(startl>>24);
cmdfis->lba4 = (uint8_t)starth;
cmdfis->lba5 = (uint8_t)(starth>>8);

cmdfis->countl = count & 0xFF;
cmdfis->counth = (count >> 8) & 0xFF;

// The below loop waits until the port is no longer busy before issuing a new command
while ((port->tfd & (ATA_DEV_BUSY | ATA_DEV_DRQ)) && spin < 1000000)
{
    spin++;
}
if (spin == 1000000)
{
    trace_ahci("Port is hung\n");
    return FALSE;
}

port->ci = 1<<slot; // Issue command

// Wait for completion
while (1)
{
    // In some longer duration reads, it may be helpful to spin on the DPS bit
    // in the PxIS port field as well (1 << 5)
    if ((port->ci & (1<<slot)) == 0)
        break;
    if (port->is & HBA_PxIS_TFES) // Task file error
    {
        trace_ahci("Read disk error\n");
        return FALSE;
    }
}

// Check again
if (port->is & HBA_PxIS_TFES)
{
    trace_ahci("Read disk error\n");
    return FALSE;
}

return true;
}

// Find a free command list slot
int find_cmdslot(HBA_PORT *port)
{
    // If not set in SACT and CI, the slot is free
    uint32_t slots = (port->sact | port->ci);
    for (int i=0; i<cmdslots; i++)
    {
        if ((slots&1) == 0)
            return i;
        slots >>= 1;
    }
    trace_ahci("Cannot find free command list entry\n");
    return -1;
}

```

Checklist

Initialisation

- Enable interrupts, DMA, and memory space access in the PCI command register
- Memory map BAR 5 register as uncacheable.
- Perform BIOS/OS handoff (if the bit in the extended capabilities is set)
- Reset controller
- Register IRQ handler, using interrupt line given in the PCI register. This interrupt line may be shared with other devices, so the usual implications of this apply.
- Enable AHCI mode and interrupts in global host control register.
- Read capabilities registers. Check 64-bit DMA is supported if you need it.
- For all the implemented ports:
 - Allocate physical memory for its command list, the received FIS, and its command tables. Make sure the command tables are 128 byte aligned.
 - Memory map these as uncacheable.
 - Set command list and received FIS address registers (and upper registers, if supported).
 - Setup command list entries to point to the corresponding command table.
 - Reset the port.
 - Start command list processing with the port's command register.
 - Enable interrupts for the port. The D2H bit will signal completed commands.
 - Read signature/status of the port to see if it connected to a drive.
 - Send IDENTIFY ATA command to connected drives. Get their sector size and count.

Start read/write command

- Select an available command slot to use.
- Setup command FIS.
- Setup PRDT.
- Setup command list entry.
- Issue the command, and record separately that you have issued it.

IRQ handler

- Check global interrupt status. Write back its value. For all the ports that have a corresponding set bit...
- Check the port interrupt status. Write back its value. If zero, continue to the next port.
- If error bit set, reset port/retry commands as necessary.
- Compare issued commands register to the commands you have recorded as issuing. For any bits where a command was issued but is no longer running, this means that the command has completed.
- Once done, continue checking if any other devices sharing the IRQ also need servicing.

External Links

- [Serial ATA Advance Host Controller Interface \(AHCI\) 1.3](http://www.intel.com/technology/serialata/ahci.htm) (<http://www.intel.com/technology/serialata/ahci.htm>)
 - [Serial ATA Revision 3.0](http://www.sata-io.org) (<http://www.sata-io.org>)
 - [ATA8-ACS, ATA8-AAM](http://www.t13.org) (<http://www.t13.org>)
 - [Haiku's AHCI implementation](https://github.com/haiku/haiku/tree/master/src/add-ons/kernel/busses/scsi/ahci) (<https://github.com/haiku/haiku/tree/master/src/add-ons/kernel/busses/scsi/ahci>)
 - [xOS AHCI implementation \(assembly language\)](https://github.com/omarrx024/xos/blob/master/kernel/blkdev/ahci.asm) (<https://github.com/omarrx024/xos/blob/master/kernel/blkdev/ahci.asm>)
-

Retrieved from "<https://wiki.osdev.org/index.php?title=AHCI&oldid=29504>"

This page was last edited on 29 April 2025, at 14:21.

This page has been accessed 42,362 times.